

1) Explain your choice of processing framework briefly.

We've chosen to do the exercises in PySpark. Python libraries have nice plotting features and jupyter notebook is a tool of choice, when it comes to combining ad-hoc analysis and reporting.

However, there are some considerations to do, when using the python API for Spark.

- Python is dynamically typed, which means that... as opposed to other API's such as Scala or Java which...?
- Parsing objects between Spark and Python (serializing) is a quite expensive task, so this should be done with care. Why ?
- Working with RDD's vs Spark Dataframes in python.... why are RDD's not available

TODO: consider when to "persist" dataframes and what it means

TODO: is there a difference between "airline" and "carrier" ?

```
In [1]: #import pyspark as spark  
#import findspark  
#import pyspark  
from pyspark.sql import SparkSession  
from pyspark.sql.functions import *  
import pandas as pd
```

```
In [5]: import matplotlib.pyplot as plt
```

Spark is set up with... Memory settings... Local xecuters (8 cores)...

```
In [2]: # https://stackoverflow.com/questions/26562033/how-to-set-apache-spark-executor  
-memory  
sc._conf.get('spark.driver.memory')
```

```
Out[2]: '30g'
```

Creating a schema for the data using af Struct type

```
In [3]: from pyspark.sql.types import StructType, StructField, IntegerType, StringType
schema = StructType([
    StructField("Year", IntegerType(), True),
    StructField("Month", IntegerType(), True),
    StructField("DayOfMonth", IntegerType(), True),
    StructField("DayOfWeek", IntegerType(), True),
    StructField("DepTime", IntegerType(), True),
    StructField("CRSDepTime", IntegerType(), True),
    StructField("ArrTime", IntegerType(), True),
    StructField("CRSArrTime", IntegerType(), True),
    StructField("UniqueCarrier", StringType(), True),
    StructField("FlightNum", IntegerType(), True),
    StructField("TailNum", StringType(), True),
    StructField("ActualElapsedTime", IntegerType(), True),
    StructField("CRSElapsedTime", IntegerType(), True),
    StructField("AirTime", IntegerType(), True),
    StructField("ArrDelay", IntegerType(), True),
    StructField("DepDelay", IntegerType(), True),
    StructField("Origin", StringType(), True),
    StructField("Dest", StringType(), True),
    StructField("Distance", IntegerType(), True),
    StructField("TaxiIn", IntegerType(), True),
    StructField("TaxiOut", IntegerType(), True),
    StructField("Cancelled", IntegerType(), True),
    StructField("CancellationCode", StringType(), True),
    StructField("Diverted", IntegerType(), True),
    StructField("CarrierDelay", IntegerType(), True),
    StructField("WeatherDelay", IntegerType(), True),
    StructField("NASDelay", IntegerType(), True),
    StructField("SecurityDelay", IntegerType(), True),
    StructField("LateAircraftDelay", IntegerType(), True)])
flights = spark.read.csv("./data/2008.csv",header=True,schema=schema, nullValue='NA')
airports = spark.read.csv("./data/airports.csv",header=True,inferSchema=True, nullValue='NA')
airlines = spark.read.csv("./data/carriers.csv",header=True,inferSchema=True, nullValue='NA')
weather = spark.read.csv("./data/weather/*daily.txt",header=True,inferSchema=True, nullValue='NA')
stations = spark.read.option("delimiter", "|").option("header", "True").csv('./data/weather/*station.txt')
```

2. How many flights were there from JFK to LAX?

Finding the number of flights from JFK to LAX

```
In [7]: flights.where((col('Origin') == 'JFK') & (col('Dest') == 'LAX')).count()
```

```
Out[7]: 8078
```

3. What was the sum and average of all arrival delays for all delayed flights?

Finding the sum and average of all arrival delays for all delayed flights Average could be found using "Describe", but to include sum, we will use select

```
In [8]: flights.where(col("ArrDelay")>0).select(avg('ArrDelay'), sum('ArrDelay')).show()
```

```
+-----+-----+
|      avg(ArrDelay) | sum(ArrDelay) |
+-----+-----+
| 32.170706265203876 |      95852748 |
+-----+-----+
```

4. What was the average departure delay for each state?

Finding the average departure delay for each state. To do this, we need the airport data from airports.csv. Instead of defining the schema explicitly as above, for illustration purposes, we'll just "infer" the schema, which means asking Spark to figure it out by presampling rows.

```
In [9]: airports.show(2)
```

```
+-----+-----+-----+-----+-----+-----+-----+
| iata | airport | city | state | country | lat | long |
+-----+-----+-----+-----+-----+-----+-----+
| 00M | Thigpen | Bay Springs | MS | USA | 31.95376472 | -89.23450472 |
| 00R | Livingston Municipal | Livingston | TX | USA | 30.68586111 | -95.01792778 |
+-----+-----+-----+-----+-----+-----+-----+
only showing top 2 rows
```

Now, let's join the dataframes, group the result on states and calculate the average departure-delay- To illustrate the "agg" function used with a map, we'll add the average arrival-delays as well

```
In [10]: # We'll do the join and persist, since we will use this dataframe later on as well
# Broadcast airports if possible
flightsWithAirports = flights.join(airports, flights.Origin == airports.iata)
```

```
In [11]: delays = flightsWithAirports.\
        groupBy(airports.state).\
        agg({"DepDelay": "avg", "ArrDelay": "avg"}).\
        select(col("state").alias("state"), \
        col("avg(DepDelay)").alias("avgDepDelay"), \
        col("avg(ArrDelay)").alias("avgArrDelay")).\
        sort(desc("avgDepDelay"))

delays.show()
```


	state	avgDepDelay	avgArrDelay
	NJ	18.28530315230682	17.073619219183303
	IL	16.037485162920703	13.927999295439097
	ME	12.60202895487689	12.972307692307693
	VT	11.906676449009538	13.11985294117647
	PA	11.706605875610164	8.359997157696183
	NY	11.581353889575762	10.433212329260538
	GA	11.47578943937115	10.746965986839188
	PR	10.823683322079676	8.884239061374899
	AR	10.697886119257086	9.709514325111076
	FL	10.617784856557332	8.554335060599021
	TX	10.484268832380778	8.741412350982355
	NH	10.483407140123559	7.463268777088934
	RI	10.345095558668053	7.284535521603119
	MD	10.136788700696506	7.0616724670931115
	SC	10.073743016759776	8.942515845928815
	NV	10.047854928293972	5.234664517182271
	WI	9.898691052537206	10.273451327433628
	VA	9.741461461852408	9.015987468487651
	VI	9.727703703703703	9.00453446191052
	MA	9.677755692715417	9.280603542532255

only showing top 20 rows

Quite difficult to get a sense of this result, so let's visualize it. PySpark does not have plotting capabilities per se, so we'll convert the Spark-dataframe to a pandas dataframe (requires installing python Pandas and Matplotlib libraries). Pandas has several easy-to-use plotting features, and sorting the data by descending departure delay will give us a visual sense of the correlation btw departure delay and arrival delay (state-wise):

```
In [10]: pdDelays = delays.toPandas()
```

```
In [10]: %matplotlib inline
states=pdDelays['state'].tolist()
#states
ax = pdDelays[['avgDepDelay','avgArrDelay']].plot(kind='bar', title ="V comp",
figsize=(15, 10), legend=True, fontsize=12)
ax.set_xticklabels(states)
ax.set_xlabel("State", fontsize=12)
ax.set_ylabel("Delay", fontsize=12)
plt.show()
```



5. Which airline performed the worst seen from a customer perspective ?

Analysing airlines, lets first load the carriers.csv file, that contains carrier-names instead of just codes. TODO - join for carrier name Broadcast join if possible

```
In [12]: airlines.show()
```

```
+----+-----+
|Code|      Description|
+----+-----+
| 02Q|      Titan Airways|
| 04Q|    Tradewind Aviation|
| 05Q|    Comlux Aviation, AG|
| 06Q| Master Top Linhas...|
| 07Q|    Flair Airlines Ltd.|
| 09Q|      Swift Air, LLC|
| 0BQ|              DCA|
| 0CQ| ACM AIR CHARTER GmbH|
| 0FQ| Maine Aviation Ai...|
| 0GQ| Inter Island Airw...|
| 0HQ| Polar Airlines de...|
| 0J |      JetClub AG|
| 0JQ|    Vision Airlines|
| 0KQ| Mokulele Flight S...|
| 0LQ|    Metropix UK, LLP.|
| 0MQ| Multi-Aero, Inc. ...|
| 0Q |    Flying Service N.V.|
| 16 |    PSA Airlines Inc.|
| 17 |    Piedmont Airlines|
| 1I | Sky Trek Int'l Ai...|
+----+-----+
```

only showing top 20 rows

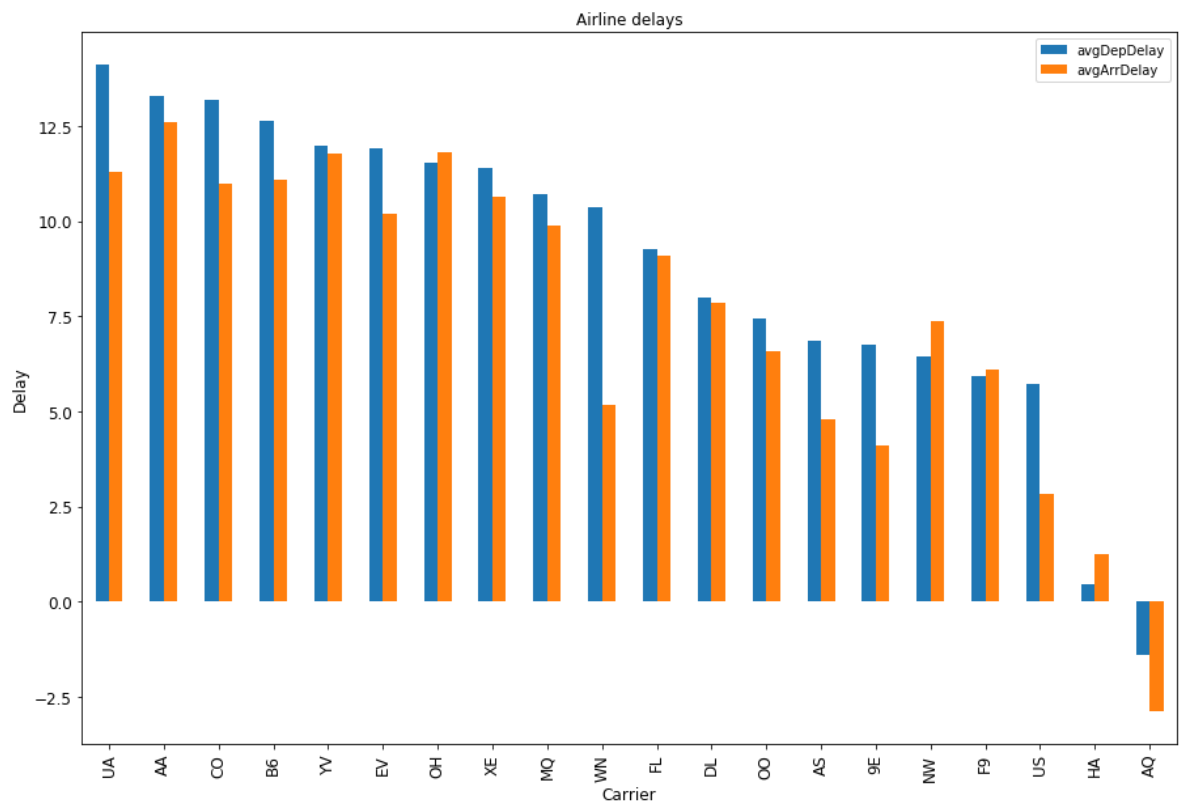
Now, lets take a look at each airline to determine, which performed the worst, see from a customer perspective. At first, lets plot the same delay-data as above, but with an airline focus instead of airport/state focus:

```
In [6]: #carrierDelays = flights.join(flights, carriers.Code == flights.UniqueCarrier).\
#      groupBy(flights.UniqueCarrier).\

carrierDelays = flights.\
    groupBy(flights.UniqueCarrier).\
        agg({"DepDelay": "avg", "ArrDelay": "avg"}).\
            select(col("UniqueCarrier").alias("UniqueCarrier"), \
                col("avg(DepDelay)").alias("avgDepDelay"), \
                col("avg(ArrDelay)").alias("avgArrDelay")).\
                    sort(desc("avgDepDelay")).toPandas()

carriers=carrierDelays['UniqueCarrier'].tolist()

ax = carrierDelays[['avgDepDelay','avgArrDelay']].plot(kind='bar', title = "Airline delays", figsize=(15, 10), legend=True, fontsize=12)
ax.set_xticklabels(carriers)
ax.set_xlabel("Carrier", fontsize=12)
ax.set_ylabel("Delay", fontsize=12)
plt.show()
print(carrierDelays.to_string(index=False))
```



UniqueCarrier	avgDepDelay	avgArrDelay
UA	14.112577	11.291322
AA	13.280898	12.607194
CO	13.185230	10.979037
B6	12.653396	11.084184
YV	12.000675	11.775181
EV	11.922538	10.208002
OH	11.536153	11.817468
XE	11.395866	10.635405
MQ	10.695642	9.890668
WN	10.383035	5.179678
FL	9.262713	9.091375
DL	8.007766	7.855163
OO	7.456443	6.598885
AS	6.848722	4.804346
9E	6.765860	4.111135
NW	6.463236	7.368539
F9	5.919602	6.108247
US	5.717490	2.848110
HA	0.455201	1.264409
AQ	-1.397783	-2.888674

Now, delays is not all that matters from a customer's point of view, so let's compute a wide range of statistics to describe the airline performance. Still, all these descriptive measures do not give us "best airline", so let's also choose a couple of them and create a general performance-measure:

- depOnTimePct
- arrOnTimePct
- completedFlightsPct

These are all percentages (eg. values between 0 and 1) describing positive features, where 1 is "perfect" and 0 is "worst". If we multiply these measures for each airline, again 1 would describe "perfect performance" and 0 would describe "worst possible performance". Let's rank the airlines according to this airline performance measure:


```

In [7]: # 1) flight-level feature engeneering
# 2) Grouping by carrier
# 3) Aggregating metrics pr. carrier
# 4) Calculation percentage metrics on carrier level

carrierPerformanceTable = flights.\
    select(flights.UniqueCarrier, \
           flights.DepDelay, \
           when(flights.DepDelay > 0,1).otherwise(0).alias("IsDepDelayed"),\
           when(flights.DepDelay > 0,0).otherwise(1).alias("IsDepOnTime"),\
           when(flights.ArrDelay > 0,1).otherwise(0).alias("IsArrDelayed"),\
           when(flights.ArrDelay > 0,0).otherwise(1).alias("IsArrOnTime"),\
           when(flights.Cancelled== 0,1).otherwise(0).alias("Completed"),\
           flights.DepDelay,
           flights.ArrDelay,
           flights.Cancelled
    ).\
    groupBy(flights.UniqueCarrier). \
    agg(sum("DepDelay").alias("DepDelay"), \
        max("DepDelay").alias("maxDepDelay"), \
        sum("ArrDelay").alias("ArrDelay"), \
        max("ArrDelay").alias("maxArrDelay"), \
        sum("IsDepDelayed").alias("isDepDelayed"), \
        sum("IsDepOnTime").alias("isDepOnTime"), \
        sum("IsArrDelayed").alias("isArrDelayed"), \
        sum("IsArrOnTime").alias("isArrOnTime"), \
        sum("Cancelled").alias("isCancelled"),\
        sum("Completed").alias("isCompleted"),\
        count(lit(1)).alias("numberOfFlights") \
    ). \
    select(col("UniqueCarrier"), \
           ((col("IsCompleted") / col("numberOfFlights"))*\
            (col("IsDepOnTime") / col("numberOfFlights"))*\
            (col("IsArrOnTime") / col("numberOfFlights"))).alias("performanceMea-
sure"),\
           round(col("IsDepOnTime") / col("numberOfFlights")*100,2).alias("depO
nTimePct"),\
           round(col("IsArrOnTime") / col("numberOfFlights")*100,2).alias("arrO
nTimePct"),\
           round(col("IsDepDelayed") / col("numberOfFlights")*100,2).alias("dep
DelayedPct"),\
           round(col("IsArrDelayed") / col("numberOfFlights")*100,2).alias("arr
DelayedPct"),\
           round(col("DepDelay") / col("isDepDelayed"),2).alias("AvgDepDelayWhe
nDelayed"),\
           round(col("ArrDelay") / col("isArrDelayed"),2).alias("AvgArrDelayWhe
nDelayed"),\
           round(col("MaxArrDelay"),2).alias("MaxArrDelay"),\
           round(col("MaxDepDelay"),2).alias("MaxDepDelay"),\
           round(col("isCancelled"),2).alias("numberOfCancelledFlights"),\
           round(col("isCancelled") / col("numberOfFlights")*100,2).alias("canc
ellationPct"),\
           round(col("isCompleted") / col("numberOfFlights")*100,2).alias("comp
letedPct"))\
    ).sort(desc("performanceMeasure")).toPandas()

carrierPerformanceTable

```

Out[7]:

	UniqueCarrier	performanceMeasure	depOnTimePct	arrOnTimePct	depDelayedPct	arrDelay
0	AQ	0.618103	82.27	75.54	17.73	24.46
1	HA	0.542419	78.55	69.70	21.45	30.30
2	9E	0.470375	73.87	65.45	26.13	34.55
3	US	0.418076	67.37	62.97	32.63	37.03
4	OO	0.411363	68.59	61.32	31.41	38.68
5	YV	0.379888	69.81	56.46	30.19	43.54
6	NW	0.370622	68.80	54.32	31.20	45.68
7	B6	0.365074	64.47	57.57	35.53	42.43
8	AS	0.364197	62.82	58.81	37.18	41.19
9	MQ	0.353339	63.21	58.07	36.79	41.93
10	OH	0.352610	71.00	51.34	29.00	48.66
11	XE	0.350882	63.71	56.58	36.29	43.42
12	FL	0.348255	63.81	55.05	36.19	44.95
13	EV	0.346459	62.72	56.25	37.28	43.75
14	DL	0.339081	64.05	53.75	35.95	46.25
15	UA	0.310971	57.48	55.40	42.52	44.60
16	WN	0.296358	49.15	60.93	50.85	39.07
17	AA	0.292314	58.43	51.52	41.57	48.48
18	F9	0.289217	56.79	51.09	43.21	48.91
19	CO	0.258553	49.84	52.53	50.16	47.47


So, all these measures still do not give us "best airline". Lets choose a couple of features:

- depOnTimePct
- arrOnTimePct
- completedFlightsPct

These are all percentages (eg. values between 0 and 1) describing positive features, where 1 is "perfect" and 0 is "worst". If we multiply these measures for each airline, again 1 would describe "perfect performance" and 0 would describe "worst possible performance". Lets rank the airlines according to this airline performance measure:

```
In [12]: carriers=carrierPerformanceTable['UniqueCarrier'].tolist()

ax = carrierPerformanceTable[['performanceMeasure']].plot(kind='bar', title ="General Airline Performance", figsize=(15, 10), legend=True, fontsize=12)
ax.set_xticklabels(carriers)
ax.set_xlabel("Carrier", fontsize=12)
ax.set_ylabel("performanceMeasure", fontsize=12)
plt.show()
```



So, it looks like this years price for best airline goes to AQ !

6. Which airport performed the worst seen from a customer perspective?

Lets do the same kind of analysis on airports, eg. which one performs worst, as seen from a customer's viewpoint. Airport performance from a customer's viewpoint could be many things. Some characteristics could also depend on whether the airport is an origination or a destination for a given flight. However, let's focus on the individual airport and ignore the origination / destination aspect.

```
In [18]: airports.show(2)
```

iata	airport	city	state	country	lat	long
00M	Thigpen	Bay Springs	MS	USA	31.95376472	-89.23450472
00R	Livingston Municipal	Livingston	TX	USA	30.68586111	-95.01792778

only showing top 2 rows

```
In [11]: # Join with airports to get destination airport name
destinationAirports=airports.select(col("iata"),col("airport")).withColumnRenamed("iata","destIata").withColumnRenamed("airport","destAirport")
# Join with airports to get origination airport info
flightsWithAirports = flights.join(destinationAirports, flights.Dest == destinationAirports.destIata).\
    alias("flightsWithDestinationAirports").\
    join(airports,flights.Origin == airports.iata)

# Pretty print the first 10 rows using pandas
flightsWithAirports.limit(10).toPandas()
```

Out[11]:

	Year	Month	DayofMonth	DayOfWeek	DepTime	CRSDepTime	ArrTime	CRSArrTime	UniqueCarrier
0	2008	1	3	4	2003	1955	2211	2225	WN
1	2008	1	3	4	754	735	1002	1000	WN
2	2008	1	3	4	628	620	804	750	WN
3	2008	1	3	4	926	930	1054	1100	WN
4	2008	1	3	4	1829	1755	1959	1925	WN
5	2008	1	3	4	1940	1915	2121	2110	WN
6	2008	1	3	4	1937	1830	2037	1940	WN
7	2008	1	3	4	1039	1040	1132	1150	WN
8	2008	1	3	4	617	615	652	650	WN
9	2008	1	3	4	1620	1620	1639	1655	WN

10 rows × 38 columns

```
In [15]: pd.set_option('display.max_columns', 250)
#flightsWithDestinationAirports.where(col("iata") == "CYS").limit(10).toPandas()
```

```
In [12]: airportPerformanceTable = flightsWithAirports.\
        select(flightsWithAirports.destIata, \
               flightsWithAirports.destAirport, \
               when(flightsWithAirports.Diverted > 0,0).otherwise(1).alias("hasArrived"),\
               ).\
        groupBy(flightsWithAirports.destIata, flightsWithAirports.destAirport). \
        agg(sum("hasArrived").alias("hasArrived"), \
            count(lit(1)).alias("numberOfFlights") \
            ). \
        select(col("destIata"), col("destAirport"), \
               col("numberOfFlights"), \
               (round(col("hasArrived") / col("numberOfFlights")*100,2)).alias("completedPct"))\
        ).sort(asc("completedPct")).limit(10).toPandas()
```

```
In [12]: airportPerformanceTable
```

```
Out[12]:
```

	destIata	destAirport	numberOfFlights	completedPct
0	OGD	Ogden-Hinckley	2	0.00
1	CYS	Cheyenne	2	0.00
2	OME	Nome	1090	95.96
3	TEX	Telluride Regional	194	96.91
4	OTZ	Ralph Wien Memorial	1086	97.42
5	WRG	Wrangell	727	97.66
6	TWF	Joslin Field - Magic Valley	1788	97.76
7	SUN	Friedman Memorial	2905	97.80
8	PSG	James C. Johnson Petersburg	727	98.07
9	HHH	Hilton Head	836	98.09

How about security delay ?

7. On appserver2 (and possibly your laptop), these files are just stored as ordinary files in the OSmanaged file system. How would they be stored in HDFS running on a cluster? Which advantages/disadvantages would that give?

The HDFS (Hadoop Distributed File System) is a distributed filesystem that supports parallelism in file reading/writing on multiple machines. This means, that every "logical" file is split into partitions, that are placed on different machines on local storage. This gives us the following benefits:

- Reading the whole logical file can be done in parallel by individual machines
- Having the partitioned data on local storage, some transformations can be performed directly on the local partition of data
- Being able to store files that are larger than any single local harddrive
- We have fault-tolerance, since all partitions are replicated three (default) times on different nodes. The partitioning scheme and replication however presents a choice between:
- Consistency, Availability and Partition tolerance

This means, that if partitioning tolerance is given in HDFS (meaning, that if one partition-replica is corrupted, the system will still be running), we need to choose between consistency and availability. HDFS offers consistency - thus, we can run into availability-issues, since a write to a file means, that to ensure that consistency, this write needs to be replicated to other replicas before being able to guarantee a consistent read of the same file. If a network (or other) failure prevents this replication, then the system is down.

Basically, this means, that the HDFS is not a high-availability system, because it gives priority to consistency.

How about security delay ?

In [13]: `airports.show(2)`

iata	airport	city	state	country	lat	long
00M	Thigpen	Bay Springs	MS	USA	31.95376472	-89.23450472
00R	Livingston Municipal	Livingston	TX	USA	30.68586111	-95.01792778

only showing top 2 rows

Clustering

```
In [21]: import urllib.request
import zipfile
import os

def downloadAndUnzip(url, filename):
    downloadFile=url+filename
    targetFile="./data/downloadStaging/"+filename
    print("Downloading and unzipping: "+downloadFile)
    urllib.request.urlretrieve(downloadFile, "./data/downloadStaging/"+filename
)

    zip_ref = zipfile.ZipFile(targetFile, 'r')
    zip_ref.extractall("./data/downloadStaging")
    zip_ref.close()
    # Cleanup
    os.system('cp ./data/downloadStaging/*daily.txt ./data/weather/')
    os.system('cp ./data/downloadStaging/*station.txt ./data/weather/')
    os.system('rm ./data/downloadStaging/*')

years=["2008"]
months=["01","02","03","04","05","06","07","08","09","10","11","12"]
for year in years:
    for month in months:
        downloadAndUnzip("https://www.ncdc.noaa.gov/orders/qclcd/"+"QCLCD"+year
+month+".zip")
```

```
Downloading and unzipping: https://www.ncdc.noaa.gov/orders/qclcd/QCLCD200801.
zip
Downloading and unzipping: https://www.ncdc.noaa.gov/orders/qclcd/QCLCD200802.
zip
Downloading and unzipping: https://www.ncdc.noaa.gov/orders/qclcd/QCLCD200803.
zip
Downloading and unzipping: https://www.ncdc.noaa.gov/orders/qclcd/QCLCD200804.
zip
Downloading and unzipping: https://www.ncdc.noaa.gov/orders/qclcd/QCLCD200805.
zip
Downloading and unzipping: https://www.ncdc.noaa.gov/orders/qclcd/QCLCD200806.
zip
Downloading and unzipping: https://www.ncdc.noaa.gov/orders/qclcd/QCLCD200807.
zip
Downloading and unzipping: https://www.ncdc.noaa.gov/orders/qclcd/QCLCD200808.
zip
Downloading and unzipping: https://www.ncdc.noaa.gov/orders/qclcd/QCLCD200809.
zip
Downloading and unzipping: https://www.ncdc.noaa.gov/orders/qclcd/QCLCD200810.
zip
Downloading and unzipping: https://www.ncdc.noaa.gov/orders/qclcd/QCLCD200811.
zip
Downloading and unzipping: https://www.ncdc.noaa.gov/orders/qclcd/QCLCD200812.
zip
```

```
In [17]: pd.set_option('display.max_columns', 250)
weatherPD = weather.limit(10).toPandas()
weatherPD
```

Out[17]:

	WBAN	YearMonthDay	Tmax	TmaxFlag	Tmin	TminFlag	Tavg	TavgFlag	Depart	DepartFlag
0	3013	20081201	51		21		36		M	
1	3013	20081202	71		21		46		M	
2	3013	20081203	51		28		40		M	
3	3013	20081204	28		14		21		M	
4	3013	20081205	39		4		22		M	
5	3013	20081206	58		16		37		M	
6	3013	20081207	67		21		44		M	
7	3013	20081208	59		32		46		M	
8	3013	20081209	32		13		23		M	
9	3013	20081210	43		9		26		M	

```
In [13]: ## pd.set_option('display.max_columns', 250)
sqlContext.registerDataFrameAsTable(stations, "stationsTable")
callSigns=sqlContext.sql("SELECT distinct WBAN as stationWBAN, CallSign from st
ationsTable").persist()
callSigns.limit(2).toPandas()
```

Out[13]:

	stationWBAN	CallSign
0	03041	MYP
1	04815	228

We'll join flight data and weather-station data to translate IATA callsign to WBAN, which is a key in weatherdata. Also, we'll construct a "yearMonthDay" column, that will be used for joining later on


```
In [55]: flightsWithStations.columns
```

```
Out[55]: ['Year',  
          'Month',  
          'DayofMonth',  
          'DayOfWeek',  
          'DepTime',  
          'CRSDepTime',  
          'ArrTime',  
          'CRSArrTime',  
          'UniqueCarrier',  
          'FlightNum',  
          'TailNum',  
          'ActualElapsedTime',  
          'CRSElapsedTime',  
          'AirTime',  
          'ArrDelay',  
          'DepDelay',  
          'Origin',  
          'Dest',  
          'Distance',  
          'TaxiIn',  
          'TaxiOut',  
          'Cancelled',  
          'CancellationCode',  
          'Diverted',  
          'CarrierDelay',  
          'WeatherDelay',  
          'NASDelay',  
          'SecurityDelay',  
          'LateAircraftDelay',  
          'stationWBAN',  
          'CallSign',  
          'yearMonthDay']
```

```
In [14]: # Start by joining flights with station data to translate IATA-code to WBAN code
flightsWithStations = flights.join(callSigns, flights.Origin==callSigns.CallSign, 'left_outer').withColumn("yearMonthDay", (concat(col('Year'), lpad(col('Month'), 2, '0'), lpad(col('DayofMonth'), 2, '0'))).cast("Integer"))
#flightsWithStations = flights.join(callSigns, flights.Dest==callSigns.CallSign, 'left_outer').withColumn("yearMonthDay", (concat(col('Year'), lpad(col('Month'), 2, '0'), lpad(col('DayofMonth'), 2, '0'))).cast("Integer"))
print("Lets check, if all flights have station information")
print("Looks like we loose flights from the following destinations: ")
flightsWithStations.where(col('stationWBAN').isNull()).groupby(col("Origin")).count().show()
#flightsWithStations.where(col('stationWBAN').isNull()).groupby(col("Dest")).count().show()
print("Lets join the weather")
flightsWithStations.limit(10).toPandas()
```

Lets check, if all flights have station information
Looks like we loose flights from the following destinations:

```
+-----+-----+
|Origin|count|
+-----+-----+
|    PSE|   755|
|    SCE|   645|
|    HHH|   836|
|    FCA|  2762|
|    CLD|  2303|
+-----+-----+
```

Lets join the weather

Out[14]:

	Year	Month	DayofMonth	DayOfWeek	DepTime	CRSDepTime	ArrTime	CRSArrTime	Unique
0	2008	1	3	4	2003	1955	2211	2225	WN
1	2008	1	3	4	754	735	1002	1000	WN
2	2008	1	3	4	628	620	804	750	WN
3	2008	1	3	4	926	930	1054	1100	WN
4	2008	1	3	4	1829	1755	1959	1925	WN
5	2008	1	3	4	1940	1915	2121	2110	WN
6	2008	1	3	4	1937	1830	2037	1940	WN
7	2008	1	3	4	1039	1040	1132	1150	WN
8	2008	1	3	4	617	615	652	650	WN
9	2008	1	3	4	1620	1620	1639	1655	WN

10 rows × 32 columns

```
In [25]: weather.count()
```

Out[25]: 359377

```
In [15]: # Now, lets join the weather information for the originating airport, using SQL
          # syntax:
          # This might be a tough one, joining 7 mill flights with 360K rows of weatherda
          # ta
          # How about the explain plan
          #sqlContext.registerDataFrameAsTable(flightsWithStations, "flightsWithStation
          #s")
          #sqlContext.registerDataFrameAsTable(weather, "weather")
          #flightsWithOriginWeather=sqlContext.sql("SELECT a.*, b.* \
          #
          #                                     from flightsWithStations a left joi
          #n \
          #
          #                                     weather
          #                                     b on (a.Ca
          #llSign = b.WBAN & \
          #
          #                                     a.yearMonthDay = b.YearMonthDay)")
          flightsWithOriginWeather=flightsWithStations.join(weather,(flightsWithStations.
          stationWBAN==weather.WBAN) & (flightsWithStations.yearMonthDay == weather.YearM
          onthDay), 'left_outer').drop('YearMonthDay')
          #flightsWithDestinationWeather=flightsWithStations.join(weather,(flightsWithSta
          tions.stationWBAN==weather.WBAN) & (flightsWithStations.yearMonthDay == weathe
          r.YearMonthDay), 'left_outer').drop('YearMonthDay')
          print("This join loses the following number of rows:")
          flightsWithOriginWeather.where(col('WBAN').isNull()).count()
```

This join loses the following number of rows:

Out[15]: 71613


From the above query, it looks like we loose

PySpark ML library - in contrast to Scikit-learn - requires features to be assembled in one column. The ML library supplies a method for doing so, vectorAssembler. Also, no NULL column are allowed, so we'll replace them with 0 (as in "no delay").

```
In [16]: dfForClustering = flightsWithOriginWeather.where(col('WBAN').isNotNull()).selec
          #t(col('UniqueCarrier'),col('DepDelay'),col('ArrDelay')).na.fill(0)
          #dfForClustering = flightsWithDestinationWeather.where(col('WBAN').isNotNull
          #()).select(col('UniqueCarrier'),col('DepDelay'),col('ArrDelay')).na.fill(0)
```

```
In [28]: dfForClustering.select(col('DepDelay'),col('ArrDelay')).toPandas().plot.scatter
          (x='DepDelay',y='ArrDelay')
```

Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc3950152b0>



Hvordan kan man være 1000 minutter forsinket i afgang men ankomme til tiden ?? Hvordan kan man være 1000 minutter forsinket i ankomst men være afgået til tiden ??

```
In [17]: dfForClustering.where(col('DepDelay').isNull()).count()
#dfForClustering.where(col('ArrDelay').isNull()).count()
#df.na.fill(0).show()
```

```
Out[17]: 0
```

```
In [71]: from pyspark.ml.stat import Correlation
dfForClustering.stat.corr("DepDelay", "ArrDelay")
```

```
Out[71]: 0.9269186899131432
```

We could cluster from a combination of departure- and arrival delay, but that would yield a measure of delay, that is not quite intuitive. As it looks like departure and arrival delays are very much correlated, it would suffice to accept arrivaldelay only as a general measure of delay. This is probably also the most important delay-type, seen from a customer view-point.

Rather than setting up low/medium/high delay-groups, we look into the flights data for hidden groups, using the Kmeans clustering method to divide the flights into "natural" delay-groups.

```
In [19]: from pyspark.ml.clustering import KMeans
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.feature import StandardScaler
%matplotlib inline
# VectorAssembler does not accept NULL values
#features = ['DepDelay']
features = ['ArrDelay']
assembler = VectorAssembler(inputCols=features, outputCol="features")
baseClusteringDF = assembler.transform(dfForClustering).cache()
baseClusteringDF.limit(10).toPandas()
```


```
Out[19]:
```

	UniqueCarrier	DepDelay	ArrDelay	features
0	WN	8	-14	[-14.0]
1	WN	19	2	[2.0]
2	WN	8	14	[14.0]
3	WN	-4	-6	[-6.0]
4	WN	34	34	[34.0]
5	WN	25	11	[11.0]
6	WN	67	57	[57.0]
7	WN	-1	-18	[-18.0]
8	WN	2	2	[2.0]
9	WN	0	-16	[-16.0]

```
In [24]: from tqdm import tqdm
elbowDict={}
numberOfClusters = range(2,10)
for cluster in tqdm(numberOfClusters):
    #print("calculating cost for k={}".format(cluster))
    kmeans = KMeans(k=cluster, seed=1) # 2 clusters here
    model = kmeans.fit(baseClusteringDF.select('features'))
    #transformed = model.transform(baseClusteringDF)
    #featuresAndPrediction = transformed.select("features", "prediction")
    WSSSE = model.computeCost(baseClusteringDF.select('features'))
    #print(str(WSSSE))
    elbowDict[cluster]=WSSSE
```

```
100%|██████████| 8/8 [01:23<00:00, 10.49s/it]
```

```
In [25]: import matplotlib.pyplot as plt
lists = sorted(elbowDict.items()) # sorted by key, return a list of tuples
x, y = zip(*lists) # unpack a list of pairs into two tuples
plt.plot(x, y, marker='*')
plt.show()
```



There is no clear "elbow" point, so we'll choose 5 clusters for the number of delay-groups

```
In [22]: cluster=5
kmeans = KMeans(k=cluster, seed=1)
model = kmeans.fit(baseClusteringDF.select('features'))
transformed = model.transform(baseClusteringDF)
```

Lets take a look at the groups:

```
In [23]: from pyspark.sql import functions as F
from pyspark.sql.window import Window
#delayGroups = transformed.groupBy("prediction").agg(avg('DepDelay').alias('avg
DepDelay'), \
#
min('DepDelay').alias
('minDepDelay'), \
#
max('DepDelay').alias
('maxDepDelay'), \
#
sum(lit(1)).alias('num
berOfFlights')).\
#withColumn('delayGroup',F.row_number().over(Window.partitionBy(lit(1)).orderBy
(col("avgDepDelay")))).cache()
#delayGroups.toPandas().sort_values(by=['delayGroup'])

delayGroups = transformed.groupBy("prediction").agg(avg('ArrDelay').alias('avgA
rrDelay'), \
min('ArrDelay').alias(
'minArrDelay'), \
max('ArrDelay').alias(
'maxArrDelay'), \
sum(lit(1)).alias('numb
erOfFlights')).\
withColumn('delayGroup',F.row_number().over(Window.partitionBy(lit(1)).orderBy(
col("avgArrDelay")))).cache()
delayGroups.toPandas().sort_values(by=['delayGroup'])
```

Out[23]:

	prediction	avgArrDelay	minArrDelay	maxArrDelay	numberOfFlights	delayGroup
0	0	-11.843752	-519	-3	3279086	1
1	2	7.662529	-2	29	2739910	2
2	4	51.709830	30	91	669026	3
3	1	132.032380	92	213	215473	4
4	3	297.201704	214	2461	34620	5

```
In [26]: flightsWithDelayGroup = flightsWithOriginWeather.where(col('WBAN').isNotNull())
.\
join(delayGroups,(coalesce(flightsWithOriginWeather.Arr
Delay,lit(0)) >=delayGroups.minArrDelay) & \
(coalesce(flightsWithOriginWeather.Arr
Delay,lit(0)) <=delayGroups.maxArrDelay)).cache()
#flightsWithDelayGroup = flightsWithDestinationWeather.where(col('WBAN').isNotN
ull()).\
#
join(delayGroups,(coalesce(flightsWithDestinationWeath
er.ArrDelay,lit(0)) >=delayGroups.minArrDelay) & \
#
(coalesce(flightsWithDestinationWeath
er.ArrDelay,lit(0)) <=delayGroups.maxArrDelay)).cache()
```

```
In [23]: #flightsWithDelayGroup.limit(10).toPandas()
# Persist to disk to be able to restart
flightsWithDelayGroup.write.mode('overwrite').parquet("./data/flightsWithDelayG
roup.parquet")
weather.write.mode('overwrite').parquet("./data/weather.parquet")
```


```
In [24]: flightsWithDelayGroup=sqlContext.read.parquet("./data/flightsWithDelayGroup.par
quet")
weather=sqlContext.read.parquet("./data/weather.parquet")
```

Lets stack the flights with delayGroup for each carrier. First we'll aggregate the carrierinfo in Spark and then use Pandas to plot.

```
In [27]: groupedUniqueCarriers = flightsWithDelayGroup.\
groupBy(col('UniqueCarrier'),col('delayGroup')).\
agg(sum(lit(1)).alias('numberOfFlights')).\
toPandas()

pt = groupedUniqueCarriers.pivot('UniqueCarrier', 'delayGroup', 'numberOfFlights')
pt.plot(kind='bar', stacked=True, figsize=(25,10))
```


Out[27]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc460239c50>



```
In [26]: groupedUniqueCarriers = flightsWithDelayGroup.\
groupBy(col('Month'),col('delayGroup')).\
agg(sum(lit(1)).alias('numberOfFlights')).\
toPandas()


pt = groupedUniqueCarriers.pivot('Month', 'delayGroup', 'numberOfFlights')
pt.plot(kind='bar', stacked=True, figsize=(25,10))
```

Out[26]: <matplotlib.axes._subplots.AxesSubplot at 0x7f5351af5ef0>



```
In [42]: transformed = model.transform(baseClusteringDF)
featuresAndPrediction = transformed.select("DepDelay", "ArrDelay", "prediction")
featuresAndPredictionPD = featuresAndPrediction.select(col('DepDelay'), col('ArrDelay'), col('prediction')).toPandas()
featuresAndPredictionPD.plot.scatter(x='DepDelay', y='ArrDelay', c='prediction', cmap='plasma')
#featuresAndPredictionPD.head(10)
```

Out[42]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc395077e10>



Classification

```
In [4]: #weather.columns
```

From <https://spark.apache.org/docs/1.6.2/ml-features.html#rformula> (<https://spark.apache.org/docs/1.6.2/ml-features.html#rformula>): RFormula selects columns specified by an R model formula. It produces a vector column of features and a double column of labels. Like when formulas are used in R for linear regression, string input columns will be one-hot encoded, and numeric columns will be cast to doubles. If not already present in the DataFrame, the output label column will be created from the specified response variable in the formula.

```
In [5]: #flightsWithDelayGroup.columns
```

Dropping "prediction" column, since this conflicts with below decisionTree, that outputs that column


```
In [ ]: #dist=flightsWithDelayGroup.agg(
dist=weather.agg(
countDistinct("Tmax"),
countDistinct("TmaxFlag"),
countDistinct("Tmin"),
countDistinct("TminFlag"),
countDistinct("Tavg"),
countDistinct("TavgFlag"),
countDistinct("Depart"),
countDistinct("DepartFlag"),
countDistinct("DewPoint"),
countDistinct("DewPointFlag"),
countDistinct("WetBulb"),
countDistinct("WetBulbFlag"),
countDistinct("Heat"),
countDistinct("HeatFlag"),
countDistinct("Cool"),
countDistinct("CoolFlag"),
countDistinct("Sunrise"),
countDistinct("SunriseFlag"),
countDistinct("Sunset"),
countDistinct("SunsetFlag"),
countDistinct("CodeSum"),
countDistinct("CodeSumFlag"),
countDistinct("Depth"),
countDistinct("DepthFlag"),
countDistinct("Water1"),
countDistinct("Water1Flag"),
countDistinct("SnowFall"),
countDistinct("SnowFallFlag"),
countDistinct("PrecipTotal"),
countDistinct("PrecipTotalFlag"),
countDistinct("StnPressure"),
countDistinct("StnPressureFlag"),
countDistinct("SeaLevel"),
countDistinct("SeaLevelFlag"),
countDistinct("ResultSpeed"),
countDistinct("ResultSpeedFlag"),
countDistinct("ResultDir"),
countDistinct("ResultDirFlag"),
countDistinct("AvgSpeed"),
countDistinct("AvgSpeedFlag"),
countDistinct("Max5Speed"),
countDistinct("Max5SpeedFlag"),
countDistinct("Max5Dir"),
countDistinct("Max5DirFlag"),
countDistinct("Max2Speed"),
countDistinct("Max2SpeedFlag"),
countDistinct("Max2Dir"),
countDistinct("Max2DirFlag")).toPandas()
dist
```

```
In [28]: flightsWithDelayGroup.limit(10).toPandas()
```

```
Out[28]:
```

	Year	Month	DayofMonth	DayOfWeek	DepTime	CRSDepTime	ArrTime	CRSArrTime	Unique
0	2008	6	29	7	1945	1945	2110	2105	WN
1	2008	6	29	7	1158	1145	1316	1305	WN
2	2008	6	29	7	1527	1530	1636	1650	WN
3	2008	6	29	7	754	800	900	920	WN
4	2008	6	29	7	954	1000	1157	1220	WN
5	2008	6	29	7	1832	1815	2042	2035	WN
6	2008	6	29	7	1639	1635	1933	1945	WN
7	2008	6	29	7	1031	1035	1338	1345	WN
8	2008	6	29	7	1113	1115	1429	1445	WN
9	2008	6	29	7	1759	1655	2119	2025	WN

10 rows × 86 columns

```
In [29]: # We'll create a binary classification target (delayedStatus)
# We'll consider delayGroup 1 and 2 as no delay, since it is such a small delay
(<29 mins)
flightsWithDelayStatus=flightsWithDelayGroup.withColumn('delayedStatus',when(flightsWithDelayGroup.delayGroup > 2, 1).otherwise(0))
flightsWithDelayStatus.persist()
```

```
Out[29]: DataFrame[Year: int, Month: int, DayofMonth: int, DayOfWeek: int, DepTime: int, CRSDepTime: int, ArrTime: int, CRSArrTime: int, UniqueCarrier: string, FlightNum: int, TailNum: string, ActualElapsedTime: int, CRSElapsedTime: int, AirTime: int, ArrDelay: int, DepDelay: int, Origin: string, Dest: string, Distance: int, TaxiIn: int, TaxiOut: int, Cancelled: int, CancellationCode: string, Diverted: int, CarrierDelay: int, WeatherDelay: int, NASDelay: int, SecurityDelay: int, LateAircraftDelay: int, stationWBAN: string, CallSign: string, WBAN: int, Tmax: string, TmaxFlag: string, Tmin: string, TminFlag: string, TavG: string, TavGFlag: string, Depart: string, DepartFlag: string, DewPoint: string, DewPointFlag: string, WetBulb: string, WetBulbFlag: string, Heat: string, HeatFlag: string, Cool: string, CoolFlag: string, Sunrise: string, SunriseFlag: string, Sunset: string, SunsetFlag: string, CodeSum: string, CodeSumFlag: string, Depth: string, DepthFlag: string, Water1: string, Water1Flag: string, SnowFall: string, SnowFallFlag: string, PrecipTotal: string, PrecipTotalFlag: string, StnPressure: string, StnPressureFlag: string, SeaLevel: string, SeaLevelFlag: string, ResultSpeed: string, ResultSpeedFlag: string, ResultDir: string, ResultDirFlag: string, AvgSpeed: string, AvgSpeedFlag: string, Max5Speed: string, Max5SpeedFlag: string, Max5Dir: string, Max5DirFlag: string, Max2Speed: string, Max2SpeedFlag: string, Max2Dir: string, Max2DirFlag: string, prediction: int, avgArrDelay: double, minArrDelay: int, maxArrDelay: int, numberOfFlights: bigint, delayGroup: int, delayedStatus: int]
```

```
In [30]: flightsWithDelayStatus.groupBy(col('delayedStatus')).count().show()
```

```
+-----+-----+
|delayedStatus| count|
+-----+-----+
|          1| 919119|
|          0|6018996|
+-----+-----+
```

We create a feature vector column for the classifier along with the label. Once this is done, we'll drop all other columns, since we do not want to carry all this data around for no use.

```
In [32]: # https://spark.apache.org/docs/2.2.0/ml-features.html#rformula
from pyspark.ml.feature import RFormula
formula = RFormula(
    #formula="delayedStatus ~ DepDelay + Tmax + TmaxFlag + Tmin + TminFlag + Ta
vg + Depart + DewPoint + WetBulb + Heat + Cool + Sunrise + Sunset + CodeSum +
    Depth + SnowFall + SnowFallFlag + PrecipTotal + PrecipTotalFlag + StnPressure
    + SeaLevel + ResultSpeed + ResultDir + AvgSpeed + Max5Speed + Max5SpeedFlag +
    Max5Dir + Max2Speed + Max2SpeedFlag + Max2Dir",
    # Lets try non weather data
    formula="delayedStatus ~ DepDelay + DepTime + Distance ",
    featuresCol="features",
    labelCol="label")

output = formula.fit(flightsWithDelayStatus.na.fill(0).na.fill('None')).transfo
rm(flightsWithDelayStatus.na.fill(0).na.fill('None')).select("features","label"
)
```

```
In [36]: output.limit(5).toPandas()
```

```
Out[36]:
```

	Year	Month	DayofMonth	DayOfWeek	DepTime	CRSDepTime	ArrTime	CRSArrTime	Unique
0	2008	6	29	7	612	615	728	735	WN
1	2008	6	29	7	1323	1255	1436	1415	WN
2	2008	6	29	7	1711	1715	1838	1835	WN
3	2008	6	29	7	930	935	1040	1055	WN
4	2008	6	29	7	1612	1555	1615	1615	WN

5 rows × 89 columns

Splitting into training- and test data. Lets count the occurrences of labels in the training data - this gives us an idea of how balanced the dataset is. We might want to balance it before training, so not to induce artificial bias towards the majority class.

```
In [33]: #output.limit(10).toPandas()
(trainingData, testData) = output.randomSplit([0.8,0.2], seed = 13234 )
trainingData.groupBy(col('Label')).count().show()
```

```
+-----+-----+
|Label|  count|
+-----+-----+
|  0.0|4815457|
|  1.0| 735037|
+-----+-----+
```

We have an unbalanced trainingset, so we'll downsample the majority class to get a balanced set. This way, we will avoid the bias in training the model. The testset however should resemble unseen data, thus we'll that unbalanced.

Reduced to 1/10th of dataset in sample fraction

```
In [34]: # Downsampling on-time flights (traininset only) to get a balanced dataset
from pyspark.sql import DataFrame
trainingDataBalanced = trainingData.where(col('label')==0).sample(False, (73503
7/4815457)/10, 42).unionAll(trainingData.where(col('label')==1).sample(False, 1
/10, 42))
trainingDataBalanced.groupBy(col('Label')).count().show()
```

```
+-----+-----+
|Label|count|
+-----+-----+
|  0.0|73165|
|  1.0|73456|
+-----+-----+
```

Saving the training set, so that we can restart the process from here later on.

```
In [35]: trainingDataBalanced.persist()
```

```
Out[35]: DataFrame[features: vector, label: double]
```

```
In [34]: #flightsWithDelayGroup.limit(10).toPandas()
# Persist to disk to be able to restart
trainingDataBalanced.write.mode('overwrite').parquet("./data/trainingDataBalanced.parquet")
testData.write.mode('overwrite').parquet("./data/testData.parquet")
```

```
In [36]: #flightsWithDelayGroup.limit(10).toPandas()
# Read disk-persistent datasets to restart
trainingDataBalanced=sqlContext.read.parquet("./data/trainingDataBalanced.parquet").persist()
testData=sqlContext.read.parquet("./data/testData.parquet").persist()
```

```
In [39]: #trainingDataBalanced.columns
# Important, when training model, because of the iterative nature
testData.columns
```

```
Out[39]: ['features', 'label']
```

```
In [32]: trainingDataBalanced.limit(5).toPandas()
```

```
Out[32]:
```

	features	label
0	(-3.0, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...	0.0
1	(-1.0, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...	0.0
2	(2.0, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...	0.0
3	(-5.0, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...	0.0
4	(-2.0, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...	0.0

Instantiate and fit the classifier - done using a pipeline.

Explain pipeline

```
In [33]: # Cross validation - too expensive
#from pyspark.ml.evaluation import BinaryClassificationEvaluator
#from pyspark.ml.classification import DecisionTreeClassifier
#from pyspark.ml import Pipeline
# Evaluate model
#evaluator = BinaryClassificationEvaluator()
# Create ParamGrid for Cross Validation
#from pyspark.ml.tuning import ParamGridBuilder, CrossValidator

#tree = DecisionTreeClassifier(labelCol="label", featuresCol="features", maxDepth=5,
#                               minInstancesPerNode=20, impurity="gini")

#paramGrid = (ParamGridBuilder()
#              .addGrid(tree.maxDepth, [1, 2, 6, 10])
#              .addGrid(tree.minInstancesPerNode, [10, 20, 40])
#              .build())

# Create 5-fold CrossValidator
#cv = CrossValidator(estimator=tree, estimatorParamMaps=paramGrid, evaluator=evaluator, numFolds=5)

# Run cross validations
#cvModel = cv.fit(trainingDataBalanced.select("features", "label"))
# Takes ~5 minutes
#print("numNodes = ", cvModel.bestModel.numNodes)
#print("depth = ", cvModel.bestModel.depth)
#print("depth = ", cvModel.bestModel.minInstancesPerNode)
```

```
In [36]: from pyspark.ml.tuning import ParamGridBuilder, TrainValidationSplit
from pyspark.ml.evaluation import BinaryClassificationEvaluator
from pyspark.ml.classification import DecisionTreeClassifier
evaluator = BinaryClassificationEvaluator()

tree = DecisionTreeClassifier(labelCol="label", featuresCol="features", impurity="gini")
paramGrid = (ParamGridBuilder()
              .addGrid(tree.maxDepth, [1, 2, 6, 10])
              .addGrid(tree.minInstancesPerNode, [10, 20, 40])
              .build())

# Create trainValidationSplit
tvS = TrainValidationSplit(estimator=tree,
                           estimatorParamMaps=paramGrid,
                           evaluator=evaluator,
                           # 80% of the data will be used for training, 20% for
                           validation.
                           trainRatio=0.8)

# Run TrainValidationSplit, and choose the best set of parameters.
model = tvS.fit(trainingDataBalanced.select("features", "label"))

# Make predictions on test data. model is the model with combination of parameters
# that performed best.
#model.transform(testData)\
#    .select("features", "label", "prediction")\
#    .show()
```

```
In [38]: treeModel = model.bestModel
treeModel
```

```
Out[38]: DecisionTreeClassificationModel (uid=DecisionTreeClassifier_4321b4fb87309d6687f4) of depth 1 with 3 nodes
```


```
In [35]: # The not so pretty Spark-way
import matplotlib.pyplot as plt
import numpy as np

from pyspark.mllib.evaluation import MulticlassMetrics
from pyspark.mllib.util import MLUtils
predictions = treeModel.transform(testData).select('prediction','label')
#predictions.show(2)
mcMetrics = MulticlassMetrics(predictions.rdd)
a= mcMetrics.confusionMatrix().toArray().transpose()
a
```

```
Out[35]: array([[1123990.,  32915.],
               [ 61306., 167014.]])
```

```
In [44]: # Pretty printing with Pandas and Matplotlib
from pyspark.mllib.evaluation import MulticlassMetrics
from pyspark.mllib.util import MLUtils
predictions = treeModel.transform(testData).select('prediction','label')

true=predictions.select('label').toPandas() #Serializing to native Python (Pandas) dataframe
predicted=predictions.select('prediction').toPandas()
pretty_print_conf_matrix(true, predicted, classes=[0,1],normalize=False,title=
'Confusion matrix',cmap=plt.cm.Blues)
```



	precision	recall	f1-score	N Obs
0.0	0.98	0.94	0.96	1203539
1.0	0.67	0.86	0.76	184082
avg	0.94	0.93	0.93	1387621

```

In [43]: import matplotlib.pyplot as plt
import itertools
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
import numpy as np

def pretty_print_conf_matrix(y_true, y_pred,
                             classes,
                             normalize=False,
                             title='Confusion matrix',
                             cmap=plt.cm.Blues):
    """
    Mostly stolen from: http://scikit-learn.org/stable/auto\_examples/model\_selection/plot\_confusion\_matrix.html#sphx-glr-auto-examples-model-selection-plot-confusion-matrix-py

    Normalization changed, classification_report stats added below plot
    """

    cm = confusion_matrix(y_true, y_pred)
    #cm = confArray

    # Configure Confusion Matrix Plot Aesthetics (no text yet)
    plt.imshow(cm, interpolation='nearest', cmap=cmap, aspect='auto')
    plt.title(title, fontsize=14)
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)
    plt.ylabel('True label', fontsize=12)
    plt.xlabel('Predicted label', fontsize=12)

    # Calculate normalized values (so all cells sum to 1) if desired
    if normalize:
        cm = np.round(cm.astype('float') / cm.sum(), 2) #(axis=1)[: , np.newaxis]

    # Place Numbers as Text on Confusion Matrix Plot
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, cm[i, j],
                 horizontalalignment="center",
                 verticalalignment = 'bottom',
                 color="white" if cm[i, j] > thresh else "black",
                 fontsize=8)

    # Add Precision, Recall, F-1 Score as Captions Below Plot
    rpt = classification_report(y_true, y_pred)
    rpt = rpt.replace('avg / total', '      avg')
    rpt = rpt.replace('support', 'N Obs')

    plt.annotate(rpt,
                 xy = (0,0),
                 xytext = (-50, -200),
                 #xytext = (0, 0),
                 xycoords='axes fraction', textcoords='offset points',
                 fontsize=12, ha='left')

    # Plot
    plt.tight_layout()

```

Lets try an SVM for predicting cancellation


```
In [40]: flightsWithDelayStatus.columns
```

```

Out[40]: ['Year',
          'Month',
          'DayofMonth',
          'DayOfWeek',
          'DepTime',
          'CRSDepTime',
          'ArrTime',
          'CRSArrTime',
          'UniqueCarrier',
          'FlightNum',
          'TailNum',
          'ActualElapsedTime',
          'CRSElapsedTime',
          'AirTime',
          'ArrDelay',
          'DepDelay',
          'Origin',
          'Dest',
          'Distance',
          'TaxiIn',
          'TaxiOut',
          'Cancelled',
          'CancellationCode',
          'Diverted',
          'CarrierDelay',
          'WeatherDelay',
          'NASDelay',
          'SecurityDelay',
          'LateAircraftDelay',
          'stationWBAN',
          'CallSign',
          'WBAN',
          'Tmax',
          'TmaxFlag',
          'Tmin',
          'TminFlag',
          'Tavg',
          'TavgFlag',
          'Depart',
          'DepartFlag',
          'DewPoint',
          'DewPointFlag',
          'WetBulb',
          'WetBulbFlag',
          'Heat',
          'HeatFlag',
          'Cool',
          'CoolFlag',
          'Sunrise',
          'SunriseFlag',
          'Sunset',
          'SunsetFlag',
          'CodeSum',
          'CodeSumFlag',
          'Depth',
          'DepthFlag',
          'Water1',
          'Water1Flag',
          'SnowFall',
          'SnowFallFlag',
          'PrecipTotal',
          'PrecipTotalFlag',
          'StnPressure',
          'StnPressureFlag',
          'SeaLevel',
          'SeaLevelFlag',
          'ResultSpeed',
          'ResultSpeedFlag',

```

```
'ResultDir',
'ResultDirFlag',
'AvgSpeed',
'AvgSpeedFlag',
'Max5Speed',
'Max5SpeedFlag',
'Max5Dir',
'Max5DirFlag',
'Max2Speed',
'Max2SpeedFlag',
'Max2Dir',
'Max2DirFlag',
'prediction',
'avgArrDelay',
'minArrDelay',
'maxArrDelay',
'numberOfFlights',
'delayGroup',
'delayedStatus']
```

```
In [45]: # https://spark.apache.org/docs/2.2.0/ml-features.html#rformula
from pyspark.ml.feature import RFormula
formula = RFormula(
    #formula="Cancelled ~ Tmax + TmaxFlag + Tmin + TminFlag + Tavg + Depart + D
ewPoint + WetBulb + Heat + Cool + Sunrise + Sunset + CodeSum + Depth + SnowFal
l + SnowFallFlag + PrecipTotal + PrecipTotalFlag + StnPressure + SeaLevel + Res
ultSpeed + ResultDir + AvgSpeed + Max5Speed + Max5SpeedFlag + Max5Dir + Max2Spe
ed + Max2SpeedFlag + Max2Dir",
    #featuresCol="features",
    formula="Cancelled ~ DepDelay + DepTime + Distance ",
    labelCol="label")

outputForSVM = formula.fit(flightsWithDelayStatus.na.fill(0).na.fill('None')).t
ransform(flightsWithDelayStatus.na.fill(0).na.fill('None')).select("features",
"label")
```

```
In [46]: #output.limit(10).toPandas()
(trainingDataForSVM, testDataForSVM) = outputForSVM.randomSplit([0.8,0.2], seed
= 13234 )
trainingDataForSVM.groupBy(col('Label')).count().show()

+-----+-----+
|Label|  count|
+-----+-----+
|  0.0|5442478|
|  1.0|108016|
+-----+-----+
```

```
In [48]: trainingDataForSVMBalanced = trainingDataForSVM.where(col('label')==0).sample(F
alse, (108016/5442478), 42).unionAll(trainingDataForSVM.where(col('label')==1))
trainingDataForSVMBalanced.groupBy(col('Label')).count().show()

+-----+-----+
|Label|  count|
+-----+-----+
|  0.0|107482|
|  1.0|108016|
+-----+-----+
```

SVM relies on distance measures. Therefore we need to normalize the data before training, to avoid that large-scaled variables are "over-weighted"

```
In [47]: trainingDataForSVMBalanced.columns
```

```
Out[47]: ['features', 'label']
```

```
In [52]: # We need to stringindex the label columns ot make the LR behave like a classifi
er
from pyspark.ml.feature import StringIndexer
indexer = StringIndexer(inputCol="label", outputCol="indexedLabel")

from pyspark.ml.feature import Normalizer
# Normalize each Vector using $L^1$ norm.
normalizer = Normalizer(inputCol="features", outputCol="normalizedFeatures", p=
2.0)
normalizedBalancedTrainingData = normalizer.transform(trainingDataForSVMBalanced).
\
    drop("features").\
    withColumnRenamed("normalizedFeatures", "features")
indexed = indexer.fit(normalizedBalancedTrainingData).transform(normalizedBalancedTrainingData).
\
    drop("features").withColumnRenamed("indexedLabel", "features")
normalizedBalancedTrainingData.limit(5).toPandas()
```

```
Out[52]:
```

	label	features
0	0.0	[-0.010499252701596863, 0.8317234096656296, 0....
1	0.0	[-0.009430697583540077, 0.6945035149021299, 0....
2	0.0	[-0.008439953257911402, 0.9350262502157561, 0....
3	0.0	[-0.0075040356671765935, 0.7087902780178619, 0...
4	0.0	[-0.017844372994792756, 0.9814405147136015, 0....

```
In [53]: from pyspark.ml.classification import LogisticRegression
lr = LogisticRegression(maxIter=10, regParam=0.3, elasticNetParam=0.8)
# Fit the model
lrModel = lr.fit(normalizedBalancedTrainingData)
```


```
In [56]: normalizer = Normalizer(inputCol="features", outputCol="normalizedFeatures", p=
2.0)
normalizedTestData = normalizer.transform(testDataForSVM).drop("features").with
ColumnRenamed("normalizedFeatures", "features")
#normalizedBalancedTestData.limit(5).toPandas()
# Predict, using the model
predictions = lrModel.transform(normalizedTestData)
predictions.groupby("prediction").count().toPandas()
```

```
Out[56]:
```

	prediction	count
0	0.0	1319270
1	1.0	68351

```
In [59]: # Pretty printing with Pandas and Matplotlib
#from pyspark.mllib.evaluation import MulticlassMetrics
#from pyspark.mllib.util import MLUtils
#predictions = treeModel.transform(testData).select('prediction','label')

true=predictions.select('label').toPandas() #Serializing to native Python (Pandas) dataframe
predicted=predictions.select('prediction').toPandas()
pretty_print_conf_matrix(true, predicted, classes=[0,1],normalize=False,title=
'Confusion matrix',cmap=plt.cm.Blues)
```



	precision	recall	f1-score	N Obs
0.0	1.00	0.97	0.98	1360370
1.0	0.40	0.99	0.57	27251
avg	0.99	0.97	0.98	1387621